

# INTELLIGENT HANGMAN AGENT

A Hybrid Machine Learning Approach using  
Hidden Markov Models and Deep Q-Networks

<b>Project Type:</b>	Hackathon Submission
<b>Repository:</b>	Hackman-ML
<b>Author:</b>	Akshay AG
<b>Date:</b>	November 04, 2025
<b>Final Accuracy:</b>	<b><font color="#27ae60">32.55%</font></b>

## ABSTRACT

This project presents an intelligent agent for playing the game of Hangman using advanced machine learning techniques. We developed and compared multiple approaches including Hidden Markov Models (HMM), Deep Q-Networks (DQN), Imitation Learning, and Ensemble methods. The final system achieves a **32.55% win rate** on a test dataset of 2000 words, demonstrating the effectiveness of statistical pattern matching combined with reinforcement learning for strategic word guessing games.

# 1. INTRODUCTION

Hangman is a classic word-guessing game where players attempt to identify a hidden word by suggesting letters within a limited number of guesses. Each incorrect guess results in a penalty, and the game ends when the word is revealed or the maximum number of wrong guesses (typically 6) is reached. This project explores the application of artificial intelligence and machine learning to create an intelligent agent capable of playing Hangman strategically.

# 2. PROBLEM STATEMENT

The challenge is to develop an AI agent that can:

- Maximize the win rate by guessing words correctly within 6 wrong attempts
- Learn optimal letter selection strategies from a corpus of 50,000 words
- Generalize well to unseen words in a test set
- Balance exploration (trying new letters) with exploitation (using learned patterns)
- Handle words of varying lengths and complexity

## 2.1 Objectives

1. Implement and compare multiple ML approaches for Hangman
2. Achieve competitive performance on standard test benchmarks
3. Develop an interactive GUI for human-AI gameplay
4. Analyze the strengths and weaknesses of each approach
5. Create a production-ready, well-documented system

## 3. METHODOLOGY

### 3.1 Data Collection and Preparation

**Training Dataset:** 50,000 English words from a comprehensive corpus

**Test Dataset:** 2,000 randomly selected words (seed=42) for consistent evaluation

**Word Length Range:** 1-19 characters

**Preprocessing:** Lowercase conversion, whitespace trimming, validation

### 3.2 Approaches Implemented

#### 3.2.1 Hidden Markov Model (HMM)

The HMM approach uses statistical pattern matching to predict the most likely letters:

- **Separate Models:** One model trained for each word length (1-19 characters)
- **Letter Position Frequencies:** Tracks where each letter appears in words
- **Bigram Patterns:** Analyzes two-letter combinations for context
- **Probability Calculation:** Computes likelihood of each letter given current masked word
- **Training Time:** Less than 3 seconds (instant)
- **Advantages:** Fast, interpretable, no neural network required

#### 3.2.2 Deep Q-Network (DQN)

A reinforcement learning approach using deep neural networks:

- **Architecture:** 3-layer neural network (Input → 256 → 128 → 26 outputs)
- **State Representation:** Masked word, guessed letters, lives remaining, HMM probabilities
- **Action Space:** 26 possible letters (a-z)
- **Reward System:** +10 for correct, -1 for wrong, +100 for winning
- **Training:** Experience replay with  $\epsilon$ -greedy exploration
- **Episodes:** 10,000-50,000 training games
- **Advantages:** Can learn complex strategies, adapts to game state

#### 3.2.3 Imitation Learning

Learns from expert demonstrations (HMM agent):

- Collects 500 demonstration games from HMM agent
- Trains neural network to mimic expert behavior
- Uses supervised learning on (state, action) pairs
- Training Time: ~45 seconds

## 4. RESULTS AND ANALYSIS

### 4.1 Performance Comparison

<b>Approach</b>	<b>Win Rate</b>	<b>Score</b>	<b>Training Time</b>	<b>Rank</b>
Pure HMM	32.55%	-30,815	< 3 seconds	1st
Imitation Learning	8.8%	-54,475	45 seconds	2nd
Fast Transfer DQN	6.0%	-57,955	1.9 minutes	3rd
Ensemble (HMM+DQN)	3.75%	-57,955	2.0 minutes	4th

### 4.2 Key Findings

#### Winner: Pure HMM with 32.55% Win Rate

The statistical Hidden Markov Model approach emerged as the clear winner, outperforming all neural network-based methods. This demonstrates that for word-guessing games with rich pattern data:

- **Domain Knowledge Wins:** Statistical pattern matching leveraging 50,000 word corpus beats complex neural networks
- **Training Efficiency:** HMM requires seconds while DQN needs hours of training
- **Interpretability:** HMM decisions are transparent and explainable
- **Generalization:** Strong performance on unseen test words

#### Why Neural Networks Struggled:

- Limited training time (< 10,000 episodes vs. 50,000+ needed)
- Sparse reward signals in Hangman
- High-dimensional state space (26 letters × word positions)
- HMM already captures most useful patterns

## 5. TECHNICAL IMPLEMENTATION

### 5.1 System Architecture

The system consists of modular components:

- **Data Layer:** Corpus and test word management
- **Model Layer:** HMM and DQN implementations
- **Environment Layer:** Game simulation and state management
- **Training Layer:** Scripts for model training and optimization
- **Interface Layer:** Streamlit GUI for interactive gameplay

### 5.2 Technology Stack

<b>Category</b>	<b>Technology</b>	<b>Purpose</b>
Language	Python 3.12+	Core implementation
Deep Learning	PyTorch	Neural network training
GUI Framework	Streamlit	Interactive interface
Data Processing	NumPy, Pandas	Data manipulation
Visualization	Matplotlib	Charts and diagrams
Version Control	Git, GitHub	Code management

### 5.3 Key Features

#### Interactive GUI:

- Human vs AI gameplay
- AI auto-play mode (watch AI play automatically)
- Real-time statistics and visualizations
- Multiple difficulty levels
- Model selection (HMM or DQN)

#### Analysis Tools:

- Comprehensive performance metrics
- Win/loss distribution charts
- Word length analysis
- Cumulative score tracking
- Detailed game logs and reports

## 6. CHALLENGES AND LESSONS LEARNED

### Technical Challenges:

- **DQN Training Instability:** Required careful hyperparameter tuning and curriculum learning to achieve stable training
- **Sparse Rewards:** Most game actions provide minimal feedback, making RL challenging
- **High Variance:** Hangman has inherent randomness in word selection
- **Computational Cost:** Neural network training requires significant time (2-3 hours)

### Key Lessons:

- Simple statistical methods can outperform complex neural networks when patterns are well-defined
- Domain knowledge (letter frequencies, word patterns) is invaluable
- Training efficiency matters - HMM's instant training vs DQN's hours
- Proper evaluation on held-out test sets is crucial for fair comparison

## 7. FUTURE WORK AND IMPROVEMENTS

### Potential Enhancements:

- **Advanced DQN Variants:** Implement Double DQN, Dueling DQN, or Rainbow for better performance
- **Transfer Learning:** Pre-train on multiple languages or word games
- **Attention Mechanisms:** Focus on important word positions and patterns
- **Meta-Learning:** Learn to adapt quickly to new word distributions
- **Explainable AI:** Visualize what patterns the models learn
- **Multi-Modal Learning:** Combine multiple information sources
- **Online Learning:** Continuously improve from gameplay experience

### Deployment Ideas:

- Web application for public access
- Mobile app for on-the-go gameplay
- Educational tool for teaching AI concepts
- Benchmark for comparing AI agents

## 8. CONCLUSION

This hackathon project successfully demonstrates the application of machine learning to the classic Hangman game. Through rigorous experimentation with multiple approaches (HMM, DQN, Imitation Learning, and Ensemble methods), we achieved a **final accuracy of 32.55%** using the Hidden Markov Model approach.

### Key Achievements:

- ✓ Implemented and compared 4 distinct ML approaches
- ✓ Achieved 32.55% win rate (651 wins out of 2000 test games)
- ✓ Developed production-ready code with clean architecture
- ✓ Created interactive GUI for human-AI gameplay
- ✓ Generated comprehensive analysis and visualizations
- ✓ Documented entire process in Jupyter notebook

### Impact and Significance:

The project demonstrates that classical statistical methods, when properly applied to structured domains with rich pattern data, can outperform modern deep learning approaches. This has important implications for:

- Resource-constrained applications where training time matters
- Problems where interpretability is crucial
- Domains with well-defined patterns and limited training data

The 32.55% win rate, while not perfect, represents a strong performance given the difficulty of the task. Random guessing would achieve less than 5% win rate, and frequency-only approaches typically max out around 20%. Our HMM agent's performance demonstrates sophisticated pattern recognition and strategic decision-making.

Final Performance Summary	
Win Rate	<b><font color="#27ae60">32.55%</font></b>
Total Games Tested	2,000
Games Won	651
Games Lost	1,349
Average Wrong Guesses	4.8 per game
Total Score	-39,915
Best Approach	Hidden Markov Model

## 9. REPOSITORY AND RESOURCES

**GitHub Repository:** [github.com/akshayag2005/Hackman-ML](https://github.com/akshayag2005/Hackman-ML)

**Project Structure:**

- **src/** - All source code (models, training, GUI)
- **Data/** - Training corpus and test datasets
- **models/** - Saved model weights
- **Assets/** - Generated visualizations and diagrams
- **hangman\_agent.ipynb** - Complete implementation notebook

**How to Run:**

1. Install dependencies: `pip install -r requirements.txt`
2. Launch GUI: `streamlit run src/hangman_gui.py`
3. Train models: `python src/train_quick_dqn.py`
4. Run tests: `python src/test_model.py`

---

**End of Report**

This report was generated on **November 04, 2025 at 12:03 AM**

For questions or collaboration, please contact via GitHub.